

Advanced Machine Learning

Emilie Chouzenoux⁽¹⁾, Frédéric Pascal⁽²⁾, Alix Chazottes⁽¹⁾ and
Houssam Zenati⁽³⁾

⁽¹⁾ Center for Computer Vision (CVN), CentraleSupélec / Opis Team, Inria

⁽²⁾ Laboratory of Signals and Systems (L2S), CentraleSupélec, University Paris-Saclay

⁽³⁾ MIND Team, Inria

{emilie.chouzenoux, alix.chazottes, frederic.pascal}@centralesupelec.fr,
houssam.zenati@inria.fr

<https://pages.saclay.inria.fr/emilie.chouzenoux/ECPml/index.htm>

<http://fredericpascal.blogspot.fr>

Dominante MDS (Mathématiques, Data Sciences)

Sept. - Nov., 2024

Contents

- 1 Introduction ML - Reminders of probability theory and mathematical statistics (Bayes, estimation, classification)
Dimension reduction and data vizualisation - FP
- 2 Robust regression approaches - EC / HZ
- 3 Stochastic approximation algorithms - EC / HZ
- 4 Hierarchical clustering - FP / AC
- 5 Mixture models fitting / Model Order Selection - FP / HZ
- 6 Nonnegative matrix factorization (NMF) - EC / HZ
- 7 Inference on graphical models - EC / HZ
- 8 Exam

Key references for this course

- Tan, P. N., Steinbach, M., Kumar V., *Data mining cluster analysis: basic concepts and algorithms. Introduction to data mining.* 2013.
- Bishop, C. M. *Pattern Recognition and Machine Learning.* Springer, 2006.
- Hastie, T., Tibshirani, R. and Friedman, J. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction.* Second edition. Springer, 2009.
- James, G., Witten, D., Hastie, T. and Tibshirani, R. *An Introduction to Statistical Learning, with Applications in R.* Springer, 2013

Course 5

(hierarchical) Clustering

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

Clustering: An Unsupervised Approach

- Extract homogeneous **meaningful** or **useful** categories from the data
- Discover/learn how the data is organized, natural structure
- No ground-truth outputs for training : **unsupervised**

Objectives

- 1 **Understanding:** Biology and medicine, finance, text mining, web, ...
- 2 **Utility:** Use cluster characteristics instead of the original data (dimension reduction, regression of high-dimensional data, ...)


The labels are unknown!

Dimension reduction vs Clustering

Let $\mathbf{X} = (x_1, \dots, x_N)$ be a set of N training samples


Dimension reduction

- Project $\mathbf{X} \in \mathbb{R}^{N,d}$ onto $\mathbf{Z} \in \mathbb{R}^{N,q}$ with $q < d$
- Visualize, denoise, reduce computational cost, ...



Clustering

- Group similar samples x_i into clusters C_k
- Based on a dissimilarity metric $\mathcal{D}(C_1, C_2)$




Clustering Applications

Market segmentation

- x : purchase history
- C_k : market segments

Medical image segmentation

- x : image pixels, voxels
- C_k : blood, muscle, tumor, ...



Text mining

- x : text, e-mails, ...
- C_k : folders, themes, ...

Key Questions on Clustering

- Types of clustering ?
- How to characterize a cluster ?
- How to define similarity or dissimilarity between samples ?
- The real/optimal number of clusters ?
- What algorithms can we use and when ?
- How to evaluate a clustering result ? (subjectivity)

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality


III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

Types of clustering: Partitional vs Hierarchical


Partitional

- Division into **non-overlapping** subsets
- Each data point is in **exactly one** subset



Hierarchical

- Clusters can have sub-clusters
- Set of nested clusters, organized as a **tree**



Types of Clusters

- **Well-separated**: Any point in a cluster is closer (or more similar) to every other point in the cluster than to any point not in the cluster.
- **Prototype-Based**: an object in a cluster is closer (more similar) to the “center” of a cluster, than to the center of any other cluster →

Assumptions about shape

- Center = **centroid** (average) or **medoid** (most representative)
- **Density-based**: dense region of points, which is separated by low-density regions, from other regions of high density. Used when the clusters are **irregular or intertwined, and when noise and outliers are present** → **Is data driven**
- **Others...** graph-based...

Distinctions between sets of clusters

- **Exclusive vs non-exclusive (overlapping)**: separate clusters vs points may belong to more than one cluster
- **Fuzzy vs non-fuzzy**: each observation $\mathbf{x}_i$ belongs to **every** cluster $\mathcal{C}_k$ with a given weight $w_k \in [0, 1]$ and $\sum_{k=1}^K w_k = 1$ (Similar to probabilistic clustering).
- **Partial vs Complete**: all data are clustered vs there may be non-clustered data, e.g., outliers, noise, “uninteresting background”...
- **Homogeneous vs Heterogeneous**: Clusters with $\neq$ size, shape, density...

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

Dissimilarity Measures

Dissimilarity is a function of the pair (x, y) : $\mathcal{D} : \mathbb{E} \times \mathbb{E} \rightarrow \mathbb{R}^+$ s.t

$$\mathcal{D}(x, y) = \mathcal{D}(y, x) \geq 0 \quad \text{and} \quad \mathcal{D}(x, x) = 0 \quad \forall x \in \mathbb{E}$$


Distance is a dissimilarity measure that satisfies also

- 1 $\mathcal{D}(x, y) = 0 \iff x = y$
- 2 $\mathcal{D}(x, y) \leq \mathcal{D}(x, z) + \mathcal{D}(z, y)$ (**metric**)


Common distances

- Minkowski: $\mathcal{D}(x, y) = \left(\sum_{j=1}^d |x_j - y_j|^q \right)^{\frac{1}{q}}$
($q = 2 \rightarrow$ Euclidian distance, $q = 1 \rightarrow$ Manhattan distance)
- Mahalanobis: $\mathcal{D}(x, y) = \left[(x - y)^T \Sigma^{-1} (x - y) \right]^{\frac{1}{2}}$
- Hamming: number of indexes where the 2 vectors differ


Minimum : $\mathcal{D}(\mathcal{C}_i, \mathcal{C}_j) = \min_{\mathbf{x} \in \mathcal{C}_i, \mathbf{y} \in \mathcal{C}_j} \mathcal{D}(\mathbf{x}, \mathbf{y})$



Maximum : $\mathcal{D}(\mathcal{C}_i, \mathcal{C}_j) = \max_{\mathbf{x} \in \mathcal{C}_i, \mathbf{y} \in \mathcal{C}_j} \mathcal{D}(\mathbf{x}, \mathbf{y})$




Group Average : $\mathcal{D}(\mathcal{C}_i, \mathcal{C}_j) = \frac{1}{n_i n_j} \sum_{\mathbf{x} \in \mathcal{C}_i} \sum_{\mathbf{y} \in \mathcal{C}_j} \mathcal{D}(\mathbf{x}, \mathbf{y})$



Between Centroids : $\mathcal{D}(\mathcal{C}_i, \mathcal{C}_j) = \mathcal{D}(m_i, m_j)$,

with $m_i = \frac{1}{n_i} \sum_{\mathbf{x} \in \mathcal{C}_i} \mathbf{x}$



Dissimilarity Between Clusters (2/2)

Objective function distances

- Ward distance: $\mathcal{D}(\mathcal{C}_i, \mathcal{C}_j) = \sqrt{\frac{2 n_i n_j}{n_i + n_j}} \mathcal{D}(m_i, m_j)$
- WPGMA (Weighted Pair Group Method with Arithmetic Mean) recursive distance

$$\mathcal{D}(\mathcal{C}_i, \mathcal{C}_j) = \frac{\mathcal{D}(\mathcal{C}_i^1, \mathcal{C}_j) + \mathcal{D}(\mathcal{C}_i^2, \mathcal{C}_j)}{2}$$

where $\mathcal{C}_i^1, \mathcal{C}_i^2$ are the child clusters of $\mathcal{C}_i$

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- **Clustering Quality**

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

What makes a good clustering ?

- Centroid: $m_i = \frac{1}{n_i} \sum_{\mathbf{x} \in \mathcal{C}_i} \mathbf{x}$
- Inertia: $J_i = \sum_{\mathbf{x} \in \mathcal{C}_i} \mathcal{D}^2(\mathbf{x}, m_i)$
(low J_i corresponds to a smaller dispersion of points around m_i .)
- Within distance: $J_w = \sum_i \sum_{\mathbf{x} \in \mathcal{C}_i} \mathcal{D}^2(\mathbf{x}, m_i) = \sum_i J_i$
- Between distance: $J_b = \sum_i n_i \mathcal{D}^2(m_i, m)$
where m is the sample mean $m = \frac{1}{n} \sum \mathbf{x}$
- Performance measures: accuracy (when ground truth is known), ARI (Adjusted Rand Index), AMI (Adjusted Mutual Information)...

A good clustering...

Minimizes the within distance J_w and maximizes the between distance J_b


Illustrative example

Objective

Cluster noisy data for a segmentation application in image processing



(a) Tree data



(b) Noisy tree data

Figure: Data on which the clustering algorithms are evaluated

Should be easy...

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

K-means

It is a **prototype-based** clustering technique.

Notations: n unlabelled data vectors of $\mathbb{R}^p$ denoted as $\mathbf{x} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$ which should be split into K classes $\mathcal{C}_1, \dots, \mathcal{C}_K$, with $\text{Card}(\mathcal{C}_k) = n_k$, $\sum_{k=1}^K n_k = n$.

Centroid of $\mathcal{C}_k$ is denoted m_k .

Optimal solution

Number of partitions of $\mathbf{x}$ into K subsets:

$$P(n, K) = \frac{1}{K!} \sum_{k=0}^K k^n (-1)^{K-k} C_K^k \text{ for } K < n$$

where $C_K^k = \frac{K!}{k! (K-k)!}$.

Example: $P(100, 5) \approx 10^{68}$!!!!

K-means algorithm

- Partitional clustering approach where K of clusters **must** be specified
- Each observation is assigned to the cluster with the closest **centroid**
- Minimizes the intra-cluster variance $V = \sum_k \sum_{i|\mathbf{x}_i \in \mathcal{C}_k} \frac{1}{n_k} \|\mathbf{x}_i - m_k\|^2$
- The basic algorithm is very simple

Algorithm 1 K-means algorithm

Input : $\mathbf{x}$ observation vectors and the number K of clusters

Output : $\mathbf{z} = (z_1, \dots, z_N)$, the labels of $(\mathbf{x}_1, \dots, \mathbf{x}_N)$

Initialization : Randomly select K points as the initial centroids

Until convergence (define a criterion, e.g. error, changes, centroids estimation...) **Repeat**

- 1 Form K clusters by assigning $\mathbf{x}_i$ to the closest centroid m_k
$$C_k = \{\mathbf{x}_i, \forall i \in \{1, \dots, n\} \mid d(\mathbf{x}_i, m_k) \leq d(\mathbf{x}_i - m_j), \forall j \in \{1, \dots, K\}\}$$
- 2 Recompute the centroids
$$\forall k \in \{1, \dots, K\} : m_k = \frac{1}{n_k} \sum_{\mathbf{x}_i \in \mathcal{C}_k} \mathbf{x}_i.$$

K-means drawbacks and alternatives







K-means is simple but ...

- Solution depends on **initialization**
- Need to **know K in advance**
- Can't handle noise or outliers : *non-robust*
- Fails with clusters of *non-convex* shapes


Several alternatives

- **K-means++**: Seeding algorithm to initialize clusters with centroids “spread-out” throughout the data
- **K-medoids**: To address the robustness aspects
- **Kernel K-means**: For overcoming the convex shape
- **Many others ...**






Correct initialization




Correct initialization




Bad initialization



Results on the data set



(a) K-means++



(b) "Clusters"

Figure: Clustering obtained with two different initialization techniques

Comments...

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN


Hierarchical clustering Principles

- Produces a set of nested clusters organized as a hierarchical tree → bypass choice of K
- Can be visualized as a **dendrogram**: a tree like diagram that records the sequences of merges or splits with **branch length corresponding to cluster distance**


Two approaches

- 1 **Agglomerative**: *Bottom-up* - Start with as much clusters as observations and iteratively *aggregate* observations using a given *distance*
- 2 **Divise**: *Top-down* - Start with one cluster containing all observations and iteratively *split* into smaller clusters

Hierarchical Clustering: The tree



(a) Hierarchical Clusters



(b) Dendrogram

We can see that ...

- Each node (cluster) in the tree (except the leaf nodes) is the union of its children (**subclusters**)
- The root of the tree is the cluster containing all objects.

Hierarchical clustering example




Figure: General principles

Inter-Cluster distance

Most popular clustering techniques

Algorithm 2 Agglomerative hierarchical clustering

Input : $\mathbf{x}$ observation vectors and “cutting” threshold λ

Output : all merged clusters set (at each iteration) and “inter-cluster” distances (between clusters)

Initialization : n = sample size = number of clusters.



While Number of clusters > 1

- 1 Compute distances between clusters
 - 2 Merged the two nearest clusters
-



Inter-Cluster distances

- MIN → Single Linkage: $d(\mathcal{C}_i, \mathcal{C}_j) = \min_{\mathbf{x} \in \mathcal{C}_i, \mathbf{y} \in \mathcal{C}_j} d(\mathbf{x}, \mathbf{y})$
- MAX → Complete Linkage: $d(\mathcal{C}_i, \mathcal{C}_j) = \max_{\mathbf{x} \in \mathcal{C}_i, \mathbf{y} \in \mathcal{C}_j} d(\mathbf{x}, \mathbf{y})$
- Group Average → Average Linkage: $d(\mathcal{C}_i, \mathcal{C}_j) = \frac{1}{n_i n_j} \sum_{\mathbf{x} \in \mathcal{C}_i} \sum_{\mathbf{y} \in \mathcal{C}_j} d(\mathbf{x}, \mathbf{y})$
- Between centroid → Centroid Linkage: $d(\mathcal{C}_i, \mathcal{C}_j) = d(m_i, m_j)$, with
$$m_i = \frac{1}{n_i} \sum_{\mathbf{x} \in \mathcal{C}_i} \mathbf{x}$$
- Objective function → Objective Linkage:
 - Ward distance $d(\mathcal{C}_i, \mathcal{C}_j) = \sqrt{\frac{2 n_i n_j}{n_i + n_j}} d(m_i, m_j)$
 - WPGMA (Weighted Pair Group Method with Arithmetic Mean)
recursive distance $d(\mathcal{C}_i, \mathcal{C}_j) = \frac{d(\mathcal{C}_i^1, \mathcal{C}_j) + d(\mathcal{C}_i^2, \mathcal{C}_j)}{2}$ where
 $\mathcal{C}_i^1, \mathcal{C}_i^2$ are the child clusters of $\mathcal{C}_i$
- ...

Different distances $\Rightarrow$ different results



(a) MIN



(b) MAX

Different distances $\Rightarrow$ different results





Figure: Group average

Ward: very similar results.

- MIN : can handle non-elliptical shape BUT sensitive to outliers, noise...
- MAX: less sensitive to outliers BUT can break large clusters and biased towards globular clusters
- Average: don't break large clusters BUT biased towards globular clusters
- Ward: Hierarchical analogue of K-means


Results on the data set - Single Linkage




(a) Noisy Tree



(b) Single Linkage




(c) Dendrogram




(d) Cutting Threshold


Results on the data set - Complete Linkage




(e) Noisy Tree



(f) Complete Linkage




(g) Dendrogram




(h) Cutting Threshold


Results on the data set - Average Linkage




(i) Noisy Tree



(j) Average Linkage




(k) Dendrogram



(l) Cutting Threshold


Results on the data set - Ward Linkage




(m) Noisy Tree



(n) Average Linkage




(o) Dendrogram




(p) Cutting Threshold


Results on the data set - WPGMA Linkage




(q) Noisy Tree



(r) Average Linkage



(s) Dendrogram



(t) Cutting Threshold

Hierarchical clustering - Pros and cons

■ Pros

- Simple and intuitive
- Unsupervised: no *a priori* assumptions
- Interpretable: number of clusters, used distance...

■ Cons

- Computational cost: single linkage ($O(n^3)$, $O(n^2)$ or $O(n)$), complete linkage ($O(n^3)$ or $O(n^2)$), average ($O(n^3)$), Ward's method ($O(n^3)$), ...
- Cutting threshold: challenging choice!
- Lack of robustness: sensitivity to outliers and noise
- No global objective function to optimize
- Handle heterogeneous data (clusters of $\neq$ size, non-globular shapes...)

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- **DBSCAN**
- HDBSCAN

DBSCAN : A Density-based Algorithm

For an observation $\mathbf{x}_i$, find a sufficiently (**MinPts**) large neighborhood (ϵ), then

- aggregate the new observations (neighbors) to the cluster $\mathcal{C}_k$ of $\mathbf{x}_i$,
- else $\mathbf{x}_i$ is an isolated observation (**outlier**).


This results in three types of points called **core**, **border**, or **noise points**.

Key parameters

- ϵ and ϵ -neighborhood: $\mathcal{N}_\epsilon(\mathbf{x}_i) = \{\mathbf{z} | d(\mathbf{x}_i, \mathbf{z}) < \epsilon\}$
- **MinPts**: n_{min} for defining core points $\mathbf{x}_i$ s.t. $\text{card}(\mathcal{N}_\epsilon(\mathbf{x}_i)) \geq n_{min}$

DBSCAN: Three Types of Points


- 1 **Core point:** is near the center of a cluster/has MinPts neighbors
- 2 **Border point:** is not a core point, but is in the neighborhood of a core point
- 3 **Noise point:** is any point that is neither a core nor a border point



DBSCAN: Influence of ϵ


The parameter ϵ represents the minimum distance between two **non-neighboring** points:

- A very large ϵ causes all possible clusters to merge into one cluster
- A very small ϵ leads to a lot of noise, points are not assigned to clusters



DBSCAN: So how do we choose ϵ ?

- Depends on the distance between the data points
- The Elbow trick on the k-NN plot is commonly used in practice (k is MinPts!):
 - x-axis all the points
 - y-axis the average distance of each point to their its k-NN
- Remains a difficult choice!



DBSCAN algorithm

Algorithm 3 DBSCAN algorithm

Input: $\mathbf{x}$ observations, ε , MinPts

Output: $\mathcal{Z}$, labels of $\mathbf{x}$

For all $\mathbf{x}_i$

- 1 Verify that $\mathbf{x}_i$ has not been visited by the algo, else $\mathbf{x}_i$ is marked “as visited”
 - 2 Identify the ε -neighborhood of $\mathbf{x}_i$, $\mathcal{N}_\varepsilon(\mathbf{x}_i)$.
 - 3 **If** $\text{card}(\mathcal{N}_\varepsilon(\mathbf{x}_i)) \leq n_{min}$, then mark P as an isolated point.
Else Create a cluster $\mathcal{C}_k$ containing $\mathbf{x}_i$ and run `class_extension( $\mathcal{C}_k, \mathbf{x}_i, \varepsilon, n_{min}$ )`
-

Cluster extension

Algorithm 4 Extension class function

Input: Cluster $\mathcal{C}_k$ to increase, observation $\mathbf{x}_i$ of $\mathcal{C}_k$, n_{min} , ε .

Output : $\mathcal{L}$ labels of observations in $\mathcal{N}_\varepsilon(\mathbf{x}_i)$

Forall $\mathbf{x}_j, i \neq j$ of $\mathcal{N}_\varepsilon(\mathbf{x}_i)$

- 1 Verify that $\mathbf{x}_j$ has not been visited by the algo, else $\mathbf{x}_i$ is marked “as visited”
 - 2 Identify the ε -neighborhood of $\mathbf{x}_j$, $\mathcal{N}_\varepsilon(\mathbf{x}_j)$.
 - 3 **If** $\text{card}(\mathcal{N}_\varepsilon(\mathbf{x}_j)) \geq n_{min}$
 $\mathcal{N}_\varepsilon(\mathbf{x}_i) = \mathcal{N}_\varepsilon(\mathbf{x}_i) + \mathcal{N}_\varepsilon(\mathbf{x}_j)$
 - 4 **If** $\mathbf{x}_j$ is not clustered, add to $\mathcal{C}_k$.
-

Illustration of DBSCAN principles

See the following video : [here](#)

Several examples can be found [here](#)





Figure: Clustering results obtained with DBSCAN algorithm.

Results on the data set - DBSCAN



(a) MinPts = 4

(b) MinPts = 256

Figure: Influence of MinPts and ϵ

Discussion: ϵ , number of clusters, MinPts...

- **Pros:** Resistant to Noise, can handle clusters of different shapes and sizes
- **Cons:** Interpretable parameters (estimation), Varying densities, High-dimensional data

Algorithms comparison




Figure: From Scikits learn: <https://ogrisel.github.io/scikit-learn.org/sklearn-tutorial/modules/clustering.html>

Today's Lecture

I. Introduction to Clustering

II. Reminders on Clustering

- Types of methods and clusters
- Distance and Dissimilarity
- Clustering Quality

III. Clustering algorithms

- K-means
- Hierarchical clustering
- DBSCAN
- HDBSCAN

HDBSCAN

Key Idea: Convert DBSCAN into a hierarchical clustering algorithm and

- bypass the choice of the ϵ -parameter!
- scan all possible solutions with all values of ϵ

Main steps:

- 1 Transform the space according to the density/sparsity
- 2 Build the minimum spanning tree of the distance weighted graph
- 3 Construct a cluster hierarchy of connected components.
- 4 Condense the cluster hierarchy based on minimum cluster size.
- 5 Extract the stable clusters from the condensed tree.

Easier to understand with an example!

Campello, R.J., Moulavi, D. and Sander, J., *"Density-based clustering based on hierarchical density estimates"*. In Pacific-Asia conference on knowledge discovery and data mining (pp. 160-172) Springer, Berlin, Heidelberg, April 2013.

HDBSCAN

Key Idea: Convert DBSCAN into a hierarchical clustering algorithm and

- bypass the choice of the ϵ -parameter!
- scan all possible solutions with all values of ϵ


Main steps:

- 1 Transform the space according to the density/sparsity
- 2 Build the minimum spanning tree of the distance weighted graph
- 3 Construct a cluster hierarchy of connected components.
- 4 Condense the cluster hierarchy based on minimum cluster size.
- 5 Extract the stable clusters from the condensed tree.

Easier to understand with an example!

Campello, R.J., Moulavi, D. and Sander, J., “*Density-based clustering based on hierarchical density estimates*”. In Pacific-Asia conference on knowledge discovery and data mining (pp. 160-172). Springer, Berlin, Heidelberg, April 2013.

HDBSCAN: Illustrative example



©https://hdbscan.readthedocs.io/en/latest/how_hdbscan_works.html

Step 1: Transform The Space

- Goal: Prepare the data for a single linkage clustering (**real data is noisy and single linkage is not robust!**)
- Key idea: Push sparse points away from the rest of the data before clustering
- The *islands/sea* analogy → Make *sea* points more distant from each other and from the *land*

How do we evaluate density ?

- Need an inexpensive density estimate $\Rightarrow$ k-NN is the simplest
- Call it the **core distance** for parameters k and point $\mathbf{x}_i$, $\text{core}_k(\mathbf{x}_i)$

And how do we connect points now ?

Step 1: Transform The Space

- Goal: Prepare the data for a single linkage clustering (**real data is noisy and single linkage is not robust!**)
- Key idea: Push sparse points away from the rest of the data before clustering
- The *islands/sea* analogy → Make *sea* points more distant from each other and from the *land*

How do we evaluate density ?

- Need an inexpensive density estimate $\Rightarrow$ k-NN is the simplest
- Call it the **core distance** for parameters k and point $\mathbf{x}_i$, $\text{core}_k(\mathbf{x}_i)$

And how do we connect points now ?



Step 1 : Mutual Reachability Distance

A new distance metric is defined as

$$d_{mreach-k}(\mathbf{x}_i, \mathbf{x}_j) = \max(\text{core}_k(\mathbf{x}_i), \text{core}_k(\mathbf{x}_j), d(\mathbf{x}_i, \mathbf{x}_j)),$$

Meaning that we want to connect points that are

- 1 Close enough to each other : $d(\mathbf{x}_i, \mathbf{x}_j)$
- 2 In a dense enough region : $\text{core}_k(\mathbf{x}_i)$




Step 2 : The Minimum Spanning Tree

- Goal: Prepare the data for clustering using d_{mreach}
- Key ideas:
 - Construct a graph that connects all points
 - Start disconnecting them by lowering a threshold (sea level drops)
 - Points are the vertices and the *edges* are weighted by d_{mreach}
 - n^2 possible edges → the minimum spanning tree


Algorithms from graph theory

- Prim's algorithm
- Dual Tree Boruvka




Step 3: Build the cluster hierarchy

Clusters emerge progressively as we lower the d_{mreach} threshold (→ sort the edges and start single linkage)




Step 4 : Condense the cluster tree

Get rid of levels that resulted in noise : $\text{nbr of points} \leq C_{\min}$
(clusters are shrinking $\neq$ splitting)




Extract the clusters

Key idea: Choose clusters that persist (live for a long time) and that are large
→ maximize a stability criterion
(flat clustering: can't select descendance of a selected cluster!)



Results



HDBSCAN: Summary

Implementation: The 5 main steps

- 1 Compute $\text{core}_k(\mathbf{x}_i)$ using *MinPts* → Measure density
- 2 Transform the space: use new metric d_{mreach}
→ Robustness to noise!
- 3 Construct a minimum spanning tree
→ Lower computational cost
- 4 Simplify/condense the tree using $C_{\min}$
→ Preprocessing for the next step
- 5 Extract final clustering results
→ Maximize cluster stability

In conclusion: Two parameters (*MinPts* and $C_{\min}$), varying densities, robust to outliers, interpretability...

HDBSCAN: Summary

Implementation: The 5 main steps

- 1 Compute $\text{core}_k(\mathbf{x}_i)$ using *MinPts* → Measure density
- 2 Transform the space: use new metric d_{mreach}
→ Robustness to noise!
- 3 Construct a minimum spanning tree
→ Lower computational cost
- 4 Simplify/condense the tree using $C_{\min}$
→ Preprocessing for the next step
- 5 Extract final clustering results
→ Maximize cluster stability

In conclusion: Two parameters (*MinPts* and $C_{\min}$), varying densities, robust to outliers, interpretability...

HDBSCAN: Summary

Implementation: The 5 main steps

- 1 Compute $\text{core}_k(\mathbf{x}_i)$ using *MinPts* → Measure density
- 2 Transform the space: use new metric d_{mreach}
→ Robustness to noise!
- 3 Construct a minimum spanning tree
→ Lower computational cost
- 4 Simplify/condense the tree using $C_{\min}$
→ Preprocessing for the next step
- 5 Extract final clustering results
→ Maximize cluster stability

In conclusion: Two parameters (*MinPts* and $C_{\min}$), varying densities, robust to outliers, interpretability...

HDBSCAN: Summary

Implementation: The 5 main steps

- 1 Compute $\text{core}_k(\mathbf{x}_i)$ using *MinPts* → Measure density
- 2 Transform the space: use new metric d_{mreach}
→ Robustness to noise!
- 3 Construct a minimum spanning tree
→ Lower computational cost
- 4 Simplify/condense the tree using $C_{\min}$
→ Preprocessing for the next step
- 5 Extract final clustering results
→ Maximize cluster stability

In conclusion: Two parameters (*MinPts* and $C_{\min}$), varying densities, robust to outliers, interpretability...

HDBSCAN: Summary

Implementation: The 5 main steps

- 1 Compute $\text{core}_k(\mathbf{x}_i)$ using *MinPts* → Measure density
- 2 Transform the space: use new metric d_{mreach}
→ Robustness to noise!
- 3 Construct a minimum spanning tree
→ Lower computational cost
- 4 Simplify/condense the tree using $C_{\min}$
→ Preprocessing for the next step
- 5 Extract final clustering results
→ Maximize cluster stability

In conclusion: Two parameters (*MinPts* and $C_{\min}$), varying densities, robust to outliers, interpretability...